THE PULSE

The Pulse: AI token spending out of control – what's next?

Details from 15 tech companies on the rapid growth of token spend, and their responses to it. Also: AI vendors can't keep up with demand, plummeting morale at Meta, and more.

GERGELY OROSZ
APR 23, 2026 · PAID

80 2 1 Share

Hello from Florida – today and tomorrow, I'm at React Miami. I've always wanted to attend this conference, and finally made it happen. If you're around, say hi!

(L-R): Myself, NeetCode founder, Navdeep Singh, & YouTuber & Twitch streamer, ThePrimeagen at React Miami

Let's get to today's topics:

1. **New trend: token spend breaks budgets – what next?** In the past 2-3 months, spending on AI agents has exploded at many tech companies, and the ramifications of this are starting to dawn on engineering leaders. We've sourced details from 15 companies, including the different ways they are coping with this realization.
2. **New trend: more AI vendors can't keep up with demand.** Related to massively increased spending, GitHub Copilot and Anthropic are starting to limit less-profitable individual users, so they can serve business users whose spend has easily 10x'd in the last few months. The exception is OpenAI and Codex.
3. **Morale at Meta hits all-time low?** Business is booming but devs at Meta are furious and worried due to looming layoffs, and an invasive tracking program rolled out to all US employees.

1. New trend: token spend breaks budgets – what next?

Last week, we covered the slightly perverse trend of "tokenmaxxing" across the industry, where devs run agents with the sole aim of boosting their personal "token stats" in an effort to rank higher on internal token leaderboards, and not be seen as a Luddite who doesn't use AI tools enough compared to peers.

This week, I spoke with a software engineer at a large company and another at a seed-stage place. Both shared almost identical stories: at their latest all-hands, company leadership expressed concerns about the fast-rising costs of tokens. At both places, token spend has increased by ~10x in the last six months – with no signs of slowing down.

I wanted to find out about this trend, so I talked to devs at 15 businesses. Below is what I learned about what's happening in workplaces of all sizes. Names are anonymized.

Large companies

Setting the default model to a cheaper one: 10,000+ person SaaS company, offices on all continents

Inside a large SaaS company, most devs use an internal background coding tool for coding. This model defaults to Claude Sonnet, which is the cheaper Claude version. Model selection is not persisted, so devs who prefer working with Opus, for instance, must reselect it on every subsequent startup.

This tool supports all major frontier models such as Sonnet, Opus, GPT, and Gemini. Devs at the company whom I talked to are very heavy users of the tool and have not encountered usage limitations.

Fintech company, US, Series D, ~8,000 people. Staff engineer:

> "The cost in token spend is off the charts – and leadership has shared this trend with us. They have not said anything beyond showing growth in spend, and mentioning that this won't be sustainable. So, nothing specific yet, but my sense is that something will have to change. Limits or prioritizing cheaper models, cutting back on hiring? Who knows."

Infra company, US, publicly traded, ~5,000 people. Engineering Director:

> **"We're monitoring but not restricting.** We are spot checking the heaviest users, but we are seeing the business cases working out.
>
> We are offering some guidance on model selection - e.g., turn off the new high-effort setting in Claude. Some users are trying open source models – but open source model usage is a bottom-up initiative, not a top-down one."

Information technology, US, 10,000+ people. Director of Engineering:

> "We have already had to raise our API budget limits multiple times in April. We recently switched to a much higher-effort level for Claude, which significantly increased the cost per PR.
>
> **One reason for the cost spike is using state-of-the-art models for demanding tasks.** We are using that high-effort setting even for fairly trivial tasks that could have been handled by much cheaper models, or even by lower-effort Claude loops. Despite a few of us pointing this out, leadership has basically said budget is not the concern right now.

> I sense that the budget increase has not been forecasted, and we're in for a reckoning. I suspect the attitude changes once finance and other cost-conscious parts of the org realize we are spending hundreds of dollars per day, per highly-engaged developer. For now, fear of missing out and not wanting to fall behind seems to be outweighing cost discipline."

Games studio, US+Europe, ~5,000 people. Senior developer:

> "What budget increase? It's very hard to get a budget for AI here! Claude Code is still not rolled out because $200/month/dev is seen as too high a cost. I talk with people at startups where $1,000/month in spending is totally normal, and it's night and day here."

Fintech company, US+Europe, late stage, ~5,000 people. Staff engineer:

> **"Some developers are now spending $500 a day (!!) on Claude Code.** Practically speaking, this means that employee costs have doubled. Productivity has increased, in my view, but now the bottleneck is code reviews. AI can spit out code quite quickly, but we still have human reviews in place. Leadership encourages using AI for code review, but my team will not blindly trust AI.
>
> The push from AI is coming from the top. This year's performance review had a section on AI, rating devs by how well they used AI, so this is another reason everyone just uses it as much as they can."

Mid-sized companies

SaaS industry, US, ~2,000 people. Dev Productivity Lead:

> "Model routing helped keep our costs growing less dramatically. For example, changing the default model reduced cost by 30%. This is our strategy with AI spend, summarized:
>
> - **Short term: spend, spend, spend!** Experiment and use whatever models make sense.
> - **Measure the impact**. Measure key outcomes and report on spend, monthly.
> - **When spend vs results diverge: adjust.** When our spend increases dramatically, but outcomes don't follow: see what we can do to adjust the

> delta. More spend should mean better outcomes. If not, we are doing something wrong."

Finance industry, US, ~2,000 people. VP of AI:

> "We have Cursor and Claude Desktop, both of which have around 800-1,200 total users. Token usage is growing somewhat unexpectedly. Estimates are being adjusted on the fly; the initial plan to have strict limits (say, $100 per user) is breaking when reality hits, and people exhaust them in 3-5 working days.
>
> Using expensive models is a problem. In regards to Cursor, many devs are defaulting to the most expensive models without realizing that going with Opus gives single percentage gains in intelligence compared to Sonnet, for example, while exhausting their budgets almost immediately.
>
> **We are working on blocking/managing out the most expensive models [with Cursor]**, as going into thousands of dollars per user, per month is not sustainable on our scale. Cursor is a good partner and we're working with them to switch to a "pooled spend" model where heavy users can tap into a pool of extra spend.
>
> Claude is a similar story. We were at $100 of Claude Desktop limit for everyone, but as we are moving forward, I can see that we would need to go much higher, especially for business-critical use cases."

Infra company, US, late-stage, ~700 people. Founder:

> "We haven't had much of an issue. Most folks police themselves for runaway costs; for example, we had someone hit like $10K in a week because they messed up caching, but it was caught and they corrected their harness.
>
> **For the most part, we don't see our high-end folks spending more than ~$1K/week.** Now, to be clear, this is not a small amount! BUT it's already a small subset of the population.
>
> We're just factoring it into engineering costs at this point: if it's, say, $2K/month per employee, that's $24K per year.
>
> Who cares, then, when engineers already cost $200-400K/year in cash comp? Okay, so what if it's $5K/month. That's $60K/year.

> **Our bet is that token costs will stabilize and we'll eventually end up with local-ish models.**
>
> Now, it could be five years before they stabilize, but overall, spend today isn't that insane to me.
>
> There's a lot of people who are just dumb about it, but most legit execs push back on this. Take the Ralph loops or other insanity where someone spends $1K/day, $5K/week or stuff like this. That's all just people being fools thinking they're doing "R&D," or somehow that they're smarter than everyone else, but they're just producing junk that never ships or is not useful.
>
> **We saw a bit of "stupid overspend" in the first couple months, but that's all gone now.** Costs could go up even more if we would "crack the whip" in wanting to see even more output, but we're not doing that."

Healthcare industry, US, ~500 people. Senior engineering manager:

> "**We are not holding back on spend, and have a monthly spend leaderboard.** And we WANT devs to spend more on tokens! For example, one of my engineers spent $1,400 on a long Claude Code session in a single day.
>
> **We are seeing massive leverage, and we do more with the same number of people.** This is why we are okay with our spending spiking. Our traffic is growing more than 10x, year-on-year, and we have managed to keep things running with the same team, and these AI tools.
>
> Engineering is now blocked on Product and Design – which never happened before! This is how fast execution has become. We now have Staff+ engineers writing Product PRDs so we can move faster.
>
> I've been in tech for close to 15 years and I never saw dramatic change like this. I just came back after a 3-month break, and every single thing is different in my day! I feel these AI agents are the biggest change in the industry since high-level languages became widespread."

E-commerce company, US & Europe, ~2,000 devs. Head of Engineering:

> "The increase in spend is INSANE. It's about usage going up, with no signs of stopping. Usage is off the charts.

> We currently do not have limits in place, and are not pausing now. Our CEO is AI-pilled and won't let us slow down.
>
> **We do buy tokens at a discount.** They start from 5% and go up with usage with the vendors we use (the usual suspects.)
>
> We don't let devs use anything lower than Opus 4.7 for coding. Cheaper models might work better, but a slight error pushed to prod would result in hours of toil."

Small companies

Series A, US, ~50 people. Principal Engineer:

> "About 15 devs are heavy users of AI and costs are rising very fast. Almost everyone uses Claude and Claude Code. We are considering four potential options:
>
> 1. **Increase AI budget, and start measuring more**. Continue doing what we are, but allow devs to use more tokens instead of hiring limits. The precise ROI is hard to quantify, but we'll start to measure and track both AI adoption and impact.
> 2. **Optimize token consumption.** Use cheaper models for simpler tasks, review token usage, and see where we can cut usage. Downside: this approach could become one with diminishing returns, fast.
> 3. **Integrate more AI providers in the company.** Find wrappers to abstract LLMs. The problem is: how do you replace Claude Code, for instance?
> 4. **Pivot to local models:** such as Kimi, Qwen, and so on. The problem is it's a big investment in high-end hardware or cloud GPUs. Upside: it offers better long-term cost control, once done.
>
> We are likely to go with option #1: increase spend BUT maintain momentum and put the right measurements in place. We can do #2, #3 and #4 later. But if we kill AI usage momentum inside the company, the outcome will probably be worse."

AI infra, US, seed stage, ~15 people. Founder:

> "**We saw a 15x increase in 6 months:**
>
> - Six months ago our spend per developer was ~$200/month
> - Today, it's around $3,000/developer/month, for our seven devs

> We're not slowing usage, especially as we are building an AI infra product. The increase was much faster than expected, though."

Small, bootstrapped company, Europe. Founding engineer:

> "Our current strategy in dealing with the increase in costs is to switch to a cheaper model; unfortunately, from Opus to Sonnet in our case. That said, Sonnet is quite decent."

How businesses manage token spend

Regardless of company size, there seems to be two strategies for how companies deal with increased spending. A summary:

Strategy #1: "let it rip and start measuring." Around half of respondents say AI spend is rising dramatically, and they have decided to do nothing about it. They *want* devs to use AI as much as it makes sense to, and to help the work as much as possible.

However, because the cost is rising dramatically, these companies are now starting to measure usage and attempting to measure the impact of their AI tools.

There's a few companies where the impact seems to be very positive, already. Smaller startups whose business is exploding in numbers of customers, load, and revenue, see that they don't need to hire more staff because existing engineers can keep supporting the growth with AI tools.

Strategy #2: curb spending. Commonly mentioned cost-saving approaches:

- Use cheaper models for simpler tasks
- Set default models to less capable ones
- Set a spending cap and make it hard for engineers to exceed it, or require consent for doing so

Most companies using strategy #1 have briefly considered going with this approach, but threw it away, because they see this approach as optimizing on the wrong thing: cutting costs before the productivity impact of using state-of-the-art tools is even known!

Discounts exist when the spend is in the millions of dollars. I asked several people if they are getting discounts from vendors when buying tokens at scale. There were no

exact numbers, but this is what I gathered in aggregate about possible custom agreements:

- **Cursor: open to discounts above a few million dollars in spend.** Companies have negotiated discounts with Cursor after crossing $1M of spending. Some companies negotiated tiered discounts from this level, starting at 5% and going higher as their spend goes up.
- **Anthropic: no discounts.** I talked with companies spending $5M+ per year on Claude which have received no discounts. If Anthropic offers discounts, it will likely be at a much higher tier.
- **All discounts are custom, so try to negotiate – it's free!** Pricing discounts are on a per-customer basis, and highly custom. The easiest way to see if a discount is available is to ask the vendors!

2. New trend: more AI vendors can't keep up with demand

Tech companies seem to be spending more on AI agents than ever before, and we see the other side of this trend play out: AI vendors struggling with high load. The actions of GitHub and Anthropic this week (see below) suggest that both are looking into ways to reduce individual customers in order to redirect compute resources to serve the business customers setting spending records.

GitHub pauses individual sign-ups & tightens usage limits

On Monday (20 April), GitHub put out this update:

1. **"New sign-ups for GitHub Copilot Pro, Pro+, and Student plans are paused.** Pausing sign-ups allows us to serve existing customers more effectively.
2. **We are tightening usage limits for individual plans.** Pro+ plans offer more than 5X the limits of Pro. Users on the Pro plan who need higher limits can upgrade to Pro+. Usage limits are now displayed in VS Code and Copilot CLI to make it easier for you to avoid hitting these limits.
3. **Opus models are no longer available in Pro plans.** Opus 4.7 remains available in Pro+ plans. As we announced in our changelog, Opus 4.5 and Opus 4.6 will be removed from Pro+."

These changes are so that GitHub can serve existing business users "with a predictable experience," which suggests GitHub is either running out of compute to serve models for business customers, or that the cost of doing so for models like Opus is rocketing and perhaps running at a loss.

Anthropic cuts Claude Code from cheaper plans – then walks it back

Anthropic caused anger this week, on Tuesday (21 April), when it silently updated its pricing page by removing Claude Code Access from the $20/month 'Pro' plan:

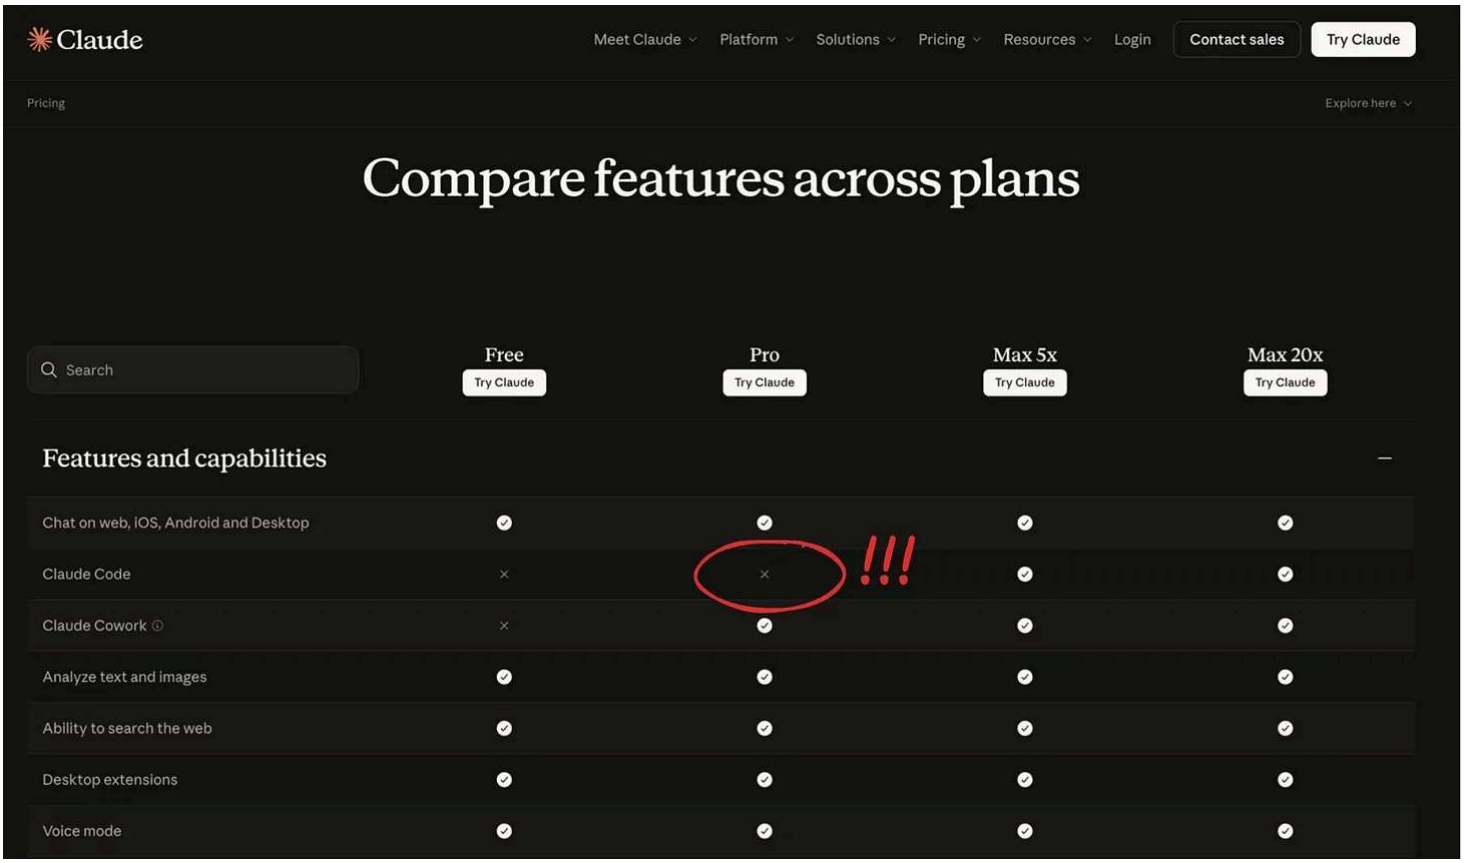

Anthropic briefly removed Claude Code access from the paid 'Pro' plan

Anthropic's Head of Growth said it was a test and they would only remove Claude Code access for 2% of new sign-ups. The pricing page was meant to "underpromise and overdeliver" in the sense of telling new customers they won't have access to Claude Code, but then still give Claude Code access to 98% of them, while leaving existing customers unaffected. This kind of experiment reminds me of the "Fake Door Test" within growth engineering, as described in our Growth Engineering deepdive with Alexey Komissarouk:

> "A Fake Door test involves offering a not-yet-available option to customers to gauge interest level. This was risky, as taking a customer who's ready to pay and

> telling them to join some kind of waiting list is a colossal waste, and risks making them feel like the target of a "bait and switch" trick.
>
> We found a way. The key insight was that people are only offended about a "bait and switch" if the "switch" is worse than the "bait." Telling customers they would pay $100 and then switching to $150 would cause a riot, but starting at $150 and then saying "just kidding, it's only $100" is a pleasant surprise."

But for Anthropic, the pricing page is assumed to be the source of truth, especially because Anthropic's image is tied to a focus upon safety and integrity! So, after the pricing page update, customers understandably assumed that all of the below were true:

1. Claude Code is no longer part of new 'Pro' plans
2. Claude Code was removed from *existing* 'Pro' plans

After all, Anthropic itself offered no additional communication on the pricing page!

The company reversed its growth experiment after a few hours – and possibly learned that things like fake door tests don't work when you are a leading AI vendor, industry-wide.

Simon Willison, independent software engineer and the creator of the the Django web framework, was non-plussed by Anthropic's unusual growth experiment, and pointed out its potential for brand damage:

> "Here's what I think happened. I think Anthropic is trying to optimize revenue growth – obviously – and someone pitched making Claude Code only available for Max and higher. That's clearly a bad idea, but "testing culture" says that it's worth putting even bad ideas out to test, just in case they surprise you.
>
> So, they started a test without taking into account the wailing and gnashing of teeth that would result when their test was noticed, or accounting for the longer-term brand damage that would be caused.
>
> Or maybe they did account for that and decided it was worth the risk.
>
> I don't think that calculation was worthwhile. They're going to have to make a very firm commitment along the lines of "we heard your feedback and we commit to

> keeping Claude Code available on our $20/month plan going forward" to regain my trust.
>
> As it stands, Codex is looking like a much safer bet for me to invest my time in learning and building educational materials around."

Anthropic's Head of Growth hinted there are capacity issues at the AI lab. Amol Avasare shared some interesting details:

> "When we launched Claude Max a year ago, it didn't include Claude Code, Cowork didn't exist, and agents that run for hours weren't a thing. Max was designed for heavy chat usage, that's it.
>
> Since then, we bundled Claude Code into Max and it took off after Opus 4. Cowork landed. Long-running async agents are now everyday workflows. The way people actually use a Claude subscription has changed fundamentally.
>
> Engagement per subscriber is way up. We've made small adjustments along the way (weekly caps, tighter limits at peak), but usage has changed a lot and our current plans weren't built for this.
>
> So, we're looking at different options to keep delivering a great experience for users. We don't know exactly what those look like yet; that's what we're testing and getting feedback on right now."

OpenAI: the only major lab not limiting usage

While both GitHub Copilot and Anthropic work on ways to reduce individual usage so they can better serve business customers, OpenAI is the sole major vendor to eschew that approach entirely. Free usage plans for Codex are alive and well, as are $20/month individual plans.

Thibault Sottiaux, Head of Codex, took a dig at Anthropic's attempt to remove Claude Code from $20/month paid plans:

> "I don't know what they are doing over there [at Anthropic], but Codex will continue to be available both in the FREE and PLUS ($20) plans. We have the compute and efficient models to support it. For important changes, we will engage with the community well ahead of making them.

> Transparency and trust are two principles we will not break, even if it means momentarily earning less. A reminder that you vote with your subscription for the values you want to see in this world."

Codex has hit 4 million active users, up from 3 million two weeks ago, and 2 million six weeks ago. Could it be that OpenAI has secured far more compute than GitHub and Anthropic?

3. Morale at Meta hits all-time low?

Meta's business is booming: the company is set to overtake Google as the world's #1 ads business by the end of the year, as covered last week. And yet, I've never encountered such poor morale inside the company after talking to half a dozen engineers there. It seems there are two major causes:

1. **Large layoffs (10%) coming.** Reuters has reported that large job cuts are due on 20 May, with around 8,000 people set to be let go. Devs at the social media giant whom I talked with are treating this report as factual. The last time Meta made large cuts was at the end of 2022 and in early 2023, when up to 25% of all staff were let go.
2. **Meta rolled out invasive software to track keystrokes and record screens for training AI.** In what sounds like something from the dystopian sci-fi show, Black Mirror, employees in the US saw a tracking tool rolled out onto their machines with no real announcement. The new tool tracks everything they do on their computer, with the goal of training Meta's new AI. The general worry folks have is that the new AI could create more layoffs, and devs fear they're training the thing that will take their jobs – with no way to opt out!

Reuters shared details about the new surveillance training software:

> "Meta is installing new tracking software on US-based employees' computers to capture mouse movements, clicks and keystrokes for use in training its artificial intelligence models, part of a broad initiative to build AI agents that can perform work tasks autonomously, the company told staffers in internal memos seen by Reuters.
>
> The tool, called Model Capability Initiative (MCI), will run on work-related apps and websites and will also take occasional snapshots of the content on employees'

> screens, according to one of the memos, posted by a staff AI research scientist on Tuesday in a channel for the company's model-building Meta SuperIntelligence Labs team."

Alex Heath [reports additional details] in his newsletter, Sources:

> "Meta's rollout of a mandatory tool that records employee keystrokes, mouse movements, and screen content to train AI sparked such intense internal backlash that an internal follow-up memo was shared Tuesday, acknowledging "there has been a lot of concern about this," according to messages I've seen.
>
> The follow-up note to Metamates (as Meta calls its workforce) attempted to offer assurances: the tool won't read files or attachments, screen content will be masked during training so the model can't memorize it, and raw data will be kept under "very tight access control."
>
> But the sign-off was less reassuring:
>
> "If you don't feel those controls are sufficient, you can control what shows up on your screen by not doing personal work on your work computer."

Engineers are upset about how little communication there has been on this tracking tool rollout. A big complaint is that Meta didn't use the company channel to announce the roll out of this tool to all devs: it was only posted in the Meta Superintelligence Labs (MSL) channel on Workspace (the internal comms tool.) That workspace only has ~25% of the company. And when an employee asked if it's possible to opt out, Meta's CTO replied that it is not: every US-based engineer will have their keystrokes logged, mouse movements captured, and work tracked. Here's what three current engineers told me about it:

Senior-level software engineer, 5+ years' tenure:

> "The situation with MCI feels dystopian. It feels to me that there's some kind of desperation and panic across the whole company. Sure, it's not unique to Meta with AI, but I don't think it's as bad elsewhere as it is here.
>
> I feel that us software engineers are no longer seen as revenue generators here, unless those of us who work on training AI models. Everyone else is seen as training data to create the next generation of models that can eventually replace our work.

> We've had layoffs at Meta before, but after the 2022 and 2023 layoffs, hiring picked up again, and we grew headcount. I have a feeling after these layoffs, headcount will not grow, so teams will have to learn to do more with less, and with AI. All while our work will be used to train the AI that can do even more of what we do."

Staff-level software engineer, 10+ years' tenure:

> "With these recent changes – pushing AI forcefully, everywhere: from token leaderboard through trajectories, to this newest tracking of everything – they managed to make building **not** be fun in this "AI-native era," when it could be so much fun to build with these tools!
>
> It feels like Meta suddenly went from the best place on earth to work, to a sad shadow of what it was."

Software engineer, 5+ years' tenure:

> "I know more than a dozen outraged devs inside the company. Also, people are sharing links to mouse jiggler software in team chats, as there's a fear that those not in front of the computer – and moving their mouse – could be targeted with the next round of layoffs.
>
> Morale is low."

Software engineer who left the company two months ago:

> "Talking to my buddies: no one feels safe at all inside the company. There are many factions building up, and everyone thinks that only the winning team will survive."
>
> Anecdotally, lots of the old-time Facebook folks (who are good) are leaving for Anthropic or OpenAI."

I've also found out that hundreds of software engineers are doing data labelling work for model training. Over the past months, several teams had 10-15% of their devs pulled off to join a team where they are doing more sophisticated data labelling and training data generation work.

It's clear enough that Mark Zuckerberg wants to build a state-of-the art coding AI model, and has decided to take unconventional actions toward this, like forcefully

enrolling US staff to generate even more high-quality training data through their daily work.

Devs never like it when they sense they are being treated as expendable, second-class citizens, and all with no notice, no say, and no way to push back against surveillance tooling and other signs that they are no longer seen as the value generators that Meta has treated them as, until now. It is a tough pill to swallow that the founder went from treating all of software engineering as a profit center, to treating it a cost center. Now, the only division treated as a profit center is the one that trains AI models.

80 Likes · 1 Restack

Discussion about this post

Comments Restacks

Write a comment...

Sri 2d

Liked by Gergely Orosz

I remember entering the industry at the start of 2023 thinking big tech was the place to be. Seeing these reports from internal Meta staff is eye opening. Can't believe they're being forced to work under surveillance, must be so dehumanising

LIKE (5) REPLY SHARE

Jakub Anderwald 4d

Liked by Gergely Orosz

I had a gut feeling that copilot cli was slower than Claude, ran a few performance tests that confirmed it - both running on opus 4.6 with the same reasoning effort showed copilot being much slower in multi-file editions, bit faster than Claude in information retrieval. Maybe that grounds Microsoft's limiting consumer access?

LIKE (1) REPLY SHARE

© 2026 Gergely Orosz · Privacy · Terms · Collection notice
Substack is the home for great culture