THE PULSE

The Pulse: Did capacity shortages turn Anthropic hostile to devs?

Also: Amazon finally allows engineers to use Claude Code and Codex, Meta forcefully assigns engineers to data labelling work ahead of layoffs, more small “AI-forward” teams, and more

GERGELY OROSZ
MAY 07, 2026 · PAID

72 4 2 Share

The Pulse is a series covering events, insights, and trends within Big Tech and startups. Notice an interesting event or trend? Hit reply and share it with me.

Today, we cover:

1. **Did Anthropic turn hostile on devs because capacity was running low?** For the past few weeks, Anthropic has continually upset devs with its “dumber” model, and by removing Claude Code access from some paid accounts. After securing lots of compute from SpaceX, could the reason have been to conceal capacity issues?
2. **Amazon finally allows Claude Code and Codex usage**. The online retail giant wanted to improve its own coding agent, Kiro, so banned other AI coding tools. But that ban is now lifted.
3. **Meta forcefully assigns engineers to data labelling ahead of job cuts.** In several teams, 20-40% of engineers are given menial, data labelling work. Could that actually boost people’s job security – for now?
4. **New trend: small “AI-forward” teams**. Meta and Amazon’s CEOs say teams of 5-10 devs do better work than teams of 50. There are important caveats: it’s unclear what they’ll do with the “excess,” and if it’s limited to “mechanical” work like rewrites.
5. **Industry Pulse**. Why Meta tracks employees’ computer activity, OpenAI starts to move off Datadog, Apple lets slip it uses Claude Code, GitHub → Xbox transfers at Microsoft, VS Code inserted “coathored by Copilot” even when Copilot did nothing, analysis of the Coinbase layoffs.

Before we start, last week I covered big pricing changes coming from GitHub, but it seems I underplayed how big they will be. Reader Julien has helpfully clarified the actual impact (thank you!)

- The multipliers will be increased for Pro and Pro+ plans *on annual renewal* (roughly a ~3x increase on average)
- With other plans (Pro, Pro+, Business, Enterprise) it's more drastic; they will adopt API token-based pricing and drop request-based pricing.

This is indeed a massive change; GitHub has heavily subsidized its usage, relative to API billing, and this change will make pricing unpredictable as of 1 June – at least initially. As Julien put it:

> "We are basically waiting to see how much we will be able to use our actual subscription with the new pricing 😇"

1. Did Anthropic turn hostile on devs because capacity was running low?

Last week, we reported on Anthropic seemingly being on a speed run to break devs' goodwill by silently "nerfing" Claude Code, banning corporate accounts without warning, and a weird growth experiment involving revoking Claude Code and then restoring it. This week, a dev on the $20/month Pro plan had Claude Code removed just days into their subscription:

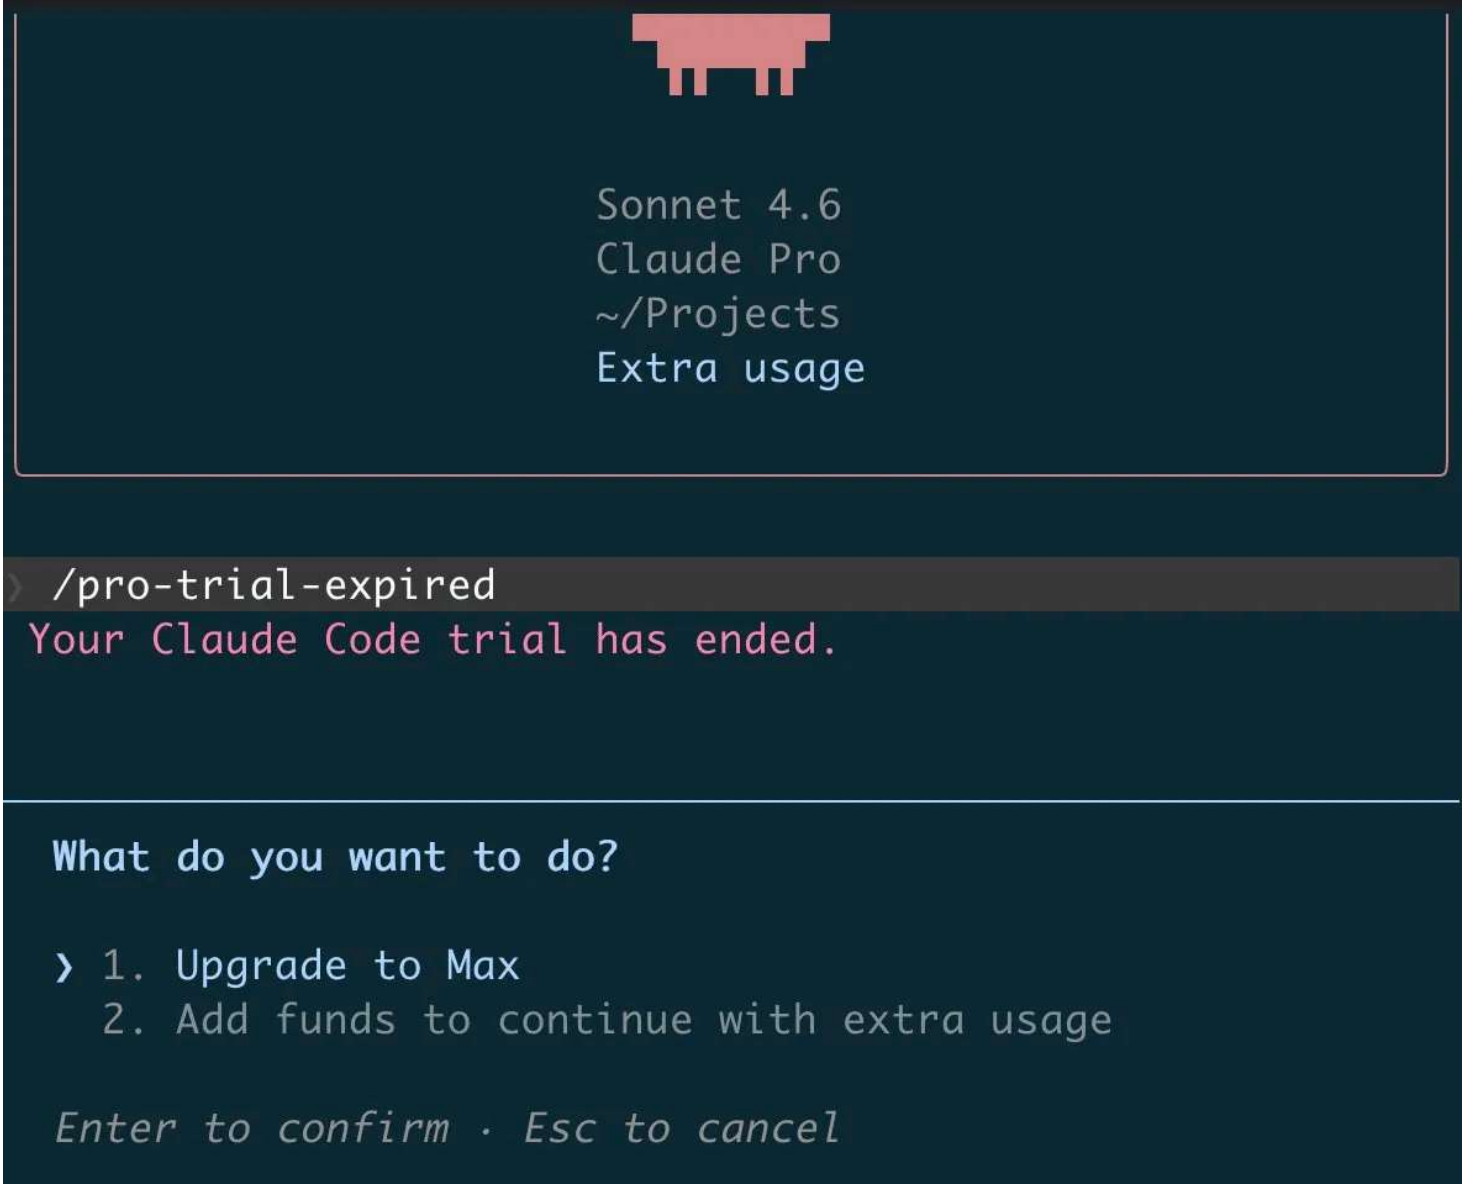

Claude Code turned out to be a trial for seven days for some paying customers. Source: Jaime Geiger

This week, Anthropic announced a big data center expansion, and relaxing previous usage limitations, while Elon Musk's SpaceX / xAI (a single company after a merger) is renting its complete Colossus 1 data center to Anthropic. From the announcement:

> "Colossus 1 features over 220,000 NVIDIA GPUs, including dense deployments of H100, H200, and next-generation GB200 accelerators. The cluster delivers extreme parallel performance for large language models, multimodal systems, scientific simulations, and generative AI at frontier scale.
>
> Anthropic plans to use this additional compute to directly improve capacity for Claude Pro and Claude Max subscribers."

In parallel with this release, Anthropic announced:

- Doubling Claude Code's current 5-hour limits for Pro, Max, Team, and seat-based Enterprise plans
- Removing peak hours limit reduction on Claude Code for Pro and Max plans
- Substantially raising API rate limits for Opus models

Is it possible that capacity issues are what led Anthropic to make Claude worse? It's confirmed the company has struggled with capacity for months. Conveniently, Claude Code being "nerfed" led to lower compute load, while removing Claude Code access from cheap plans could look like rate limiting. Even the banning of corporate accounts could be seen as scaling back at a time when the business has struggled to serve existing growth. Yesterday, (6 May), at the Code with Claude event hosted by Anthropic, CEO, Dario Amodei, said:

> "We originally planned for 10x growth, and we've seen something more like 80x growth in revenue and usage over the last period of time."

SpaceX / xAI renting a good chunk of its capacity to Anthropic is ironic, considering that xAI (Musk's AI startup) builds Grok, a frontier model and direct rival of Claude, and also in January, Anthropic banned xAI developers from Claude. As covered at the time:

> "It's common for an AI lab to not allow another AI lab to use its model, like at OpenAI, Anthropic, and Google. On the other side, there's also the pertinent question of why a leading AI lab would even want to use a rival for its own day-to-day work?
>
> Turns out, xAI (Elon Musk's AI lab) was relying on Cursor to write code, which we know because they got cut off."

Anthropic likely banned xAI to stop Claude from being potentially distilled while it tried to improve Grok's coding capability. Meanwhile, Musk called Anthropic "misanthropic and evil" earlier this year, and said the new tenant "hates Western civilization". But both parties seem happy to put that behind them and strike a deal, so perhaps there's something else at play.

Could SpaceX / xAI be checking out of the frontier-AI model wars? Leasing a good chunk of its data center capacity might suggest that. SpaceX / xAI has two data

centers: Colossus 1 and Colossus 2. Colossus 1 represents somewhere around 45% of current SpaceX / xAI capacity, and 20-25% of planned total capacity.

Giving up as much capacity as this might indicate a lack of demand, or capacity sitting idle. It also means Grok is losing out in market share to Claude, ChatGPT, and other leading models. *In February's AI tooling survey we found scarce mention of Grok, which lagged in usage behind open models like DeepSeek and Qwen.*

To be fair, unlike Anthropic and OpenAI, Grok never had a B2C nor B2B business that took off. The biggest consumer use case for Grok seems to be its integration into the social media platform, X; at least, I don't know of any tech company using the model for serious work.

"The enemy of my enemy is my friend", says the maxim, and if there's one company Musk hates, it's OpenAI. He is currently suing OpenAI, claiming it betrayed its founding nonprofit mission to develop safe AGI for humanity's benefit by shifting to a profit-driven model backed by Microsoft. Musk also claims that despite investing about $40M, he has no ownership of the company.

He wants $150B in damages, the removal of Sam Altman and Greg Brockman, and for OpenAI to return to a full nonprofit, as per when he invested in the company. *We covered more about OpenAI's own ethical challenges between nonprofit and for-profit right after the firing of Sam Altman in 2023, in the deepdive What is OpenAI, really?*

Similarly, Anthropic may well have an issue with OpenAI, if CEO Dario Amodei's failure to join hands with Sam Altman while sharing a stage with the Prime Minister of India earlier this year is anything to go by.

(Most) AI leaders join hands at the AI Impact Summit with India's Prime Minister.
Source: Fortune

Capacity issues hurting Anthropic would benefit OpenAI, and so by offering significant capacity to Anthropic, Musk is making it harder for OpenAI to win the market. That would be ironic, given he's a former investor.

2. Amazon finally allows Claude Code and Codex usage

In February, I reported that despite Amazon being an investor in Anthropic, the e-commerce giant bans using Claude Code internally. From The Pulse at the time:

> "Amazon is one of the biggest investors in Anthropic, and its devs are heavy users of Claude models. Even so, Amazon engineers are banned from using Claude Code because Amazon wants them to dogfood its own AI coding tool, Kiro.
>
> Kiro is alright, but not nearly as powerful as Claude Code. Amazon clearly wants to make it better, so chooses to lock out its devs from using a more capable tool from a company which it is heavily invested in."

Amazon engineers I talked to were disappointed, especially because Claude Code works better in many scenarios than Kiro. But this week the decision was reversed, I've learned from engineers at the company.

Amazon engineers can use Claude Code and Codex internally – on one condition: the tool needs to run via Bedrock (Amazon's managed platform for running AI models and agents).

I asked devs how they find Kiro vs Claude Code:

- Kiro seems to work better for niche, internal use cases. The model Kiro uses is trained on internal tools and systems. Oncall engineers seem to prefer Kiro as the go-to tool.
- Claude Code is often the choice for building something new and for planning new systems.

There's now "proper" competition inside Amazon, and I'd expect the Kiro team to double down on making their own, internal tool into the most-used AI agent via better integration of internal workflows.

3. Meta forcefully assigns engineers to data labelling, pre-layoff

Another detail I unearthed this week was from talking with Meta engineers, who say many teams there have had around 20-40% of their devs forcefully reassigned to data labelling jobs. There are easily 1,000+ software engineers doing data labelling at Meta now.

Data labelling is used to improve Meta's own coding model during reinforcement learning (RL). The data labelling work is roughly like this:

- Open an internal repository marked for data labelling
- Create issues you spot, and add test cases helpful, or which are missing
- Package the whole repository into a Docker container, and submit it for the training team to use

It's a pretty mundane job, and less fulfilling than building products; also, you know you're helping the model get better at software engineering and that it may one day

take your job as a result.

Understandably, there are mixed feelings about this redeployment, with layoffs coming soon. On Wednesday, 20 May, Meta will announce layoffs. Perhaps those moved to do data labelling could actually be "safer" than colleagues on product teams. *Of course, this is speculation, but it would be cruel if Meta cut devs reassigned to data labelling.*

Either way, the feeling inside the company is that devs have little agency or control. Those who have been reassigned this week start data labelling next week; if the model doesn't need as much data labelling as currently thought, then these folks would likely be moved on again. This is especially significant for Meta because its culture used to be all about engineers *having a choice*.

Previously, new devs would go through a weeks-long bootcamp and choose the teams *they* wanted to work on. That looks like it's history.

4. New trend: small "AI-forward" teams

A potential new trend could be of smaller teams becoming the norm, based on similar comments by two leading leaders:

Meta: a 10-person team does the work of a team of 50. Mark Zuckerberg, on an internal company meeting a week ago (Thursday, 30 April), said, as per Business Insider:

> "If a team used to take 50 or 100 people and now it takes 10, having 50 or 100 people on that team can actually be counterproductive going forward so I think we need to fix that."

Zuckerberg said that they plan to spin up more of these teams, and do more overall, thanks to having more capability to execute. "More" could mean mobile apps, although my guess is it will mean more of everything: more products inside the major apps and more services as well.

One possible reason why the "build more apps" strategy might not work out great: it didn't before! There was a time when Meta had lots of additional apps, which simply never gained traction. In 2014-2015, the company launched and then killed these standalone apps:

- Poke: an early Snapchat clone with self-destructing messages (launched: 2012, killed: 2014)
- Slingshot: a Snapchat-style photo messaging app (launched: 2014, killed: 2015)
- Rooms: pseudo-anonymous, topic-based chat rooms (launched: 2014, killed: 2015)
- Riff – collaborative short-form video creation (launched: 2015, killed: 2016)
- Lifestage – teens-only video profile app (launched: 2016, killed: 2017)

Perhaps Meta's bet is that they can build 50 apps instead of five in a year, and one or two will become winners. I'm somewhat skeptical, but we'll see soon enough.

Amazon: a team of five rewriting a service 5x faster than a team of 50. CEO Andy Jassy gave this example last Wednesday (29 April) on Amazon's earning call:

> "I mean just look at how coding, agentic coding is changing how we're all building products. I think it's going to have a comparable impact on how we do DevOps and how we do customer service, how we do research, how we do analytics, how sales is conducted. I think every single one of these functions that we all do at work are going to very significantly change. (...)
>
> If you look at one of our services, we swapped out the engine of the service while we are also running the service, full-tilt. Normally, that would have taken 40 or 50 people about a year to do. We took 5, really smart, AI forward-thinking people, [used]agentic coding tools, and those 5 people rebuilt [the service] in 65 days. That is a very different world of operating, and that's the world I think we're heading to over the next few years."

While this is impressive, let's not forget that it's not about building a new thing, but rewriting an existing service; one that was likely covered with tests and had well-understood use cases. LLMs work great for migrations and rewrites, and when the spec is clear and the end result easy to test, they are a great choice! But LLMs don't help much to build something new, or when the requirements are unclear, or there's no prior art to copy.

Even so, the trend looks real; all the startups I talk with find success from working with smaller teams, experienced engineers, and AI tools. Talking with engineers in Big Tech: they say that teams seem 20-30% smaller than two years ago, but get about the same amount of work done, with AI tools filling the gap.

5. Industry Pulse

Why Meta tracks employees' computer activity

Mark Zuckerberg explained to devs why every part of their work is being tracked. The Information reports on an internal meeting:

> "Mark Zuckerberg said using employees' computer activity to train its AI models could give the company an edge over rivals, arguing that its staff are of higher average intelligence than contractors typically used by data-labeling firms."

I personally don't buy this flattering argument that Meta gets an edge because its own people are smarter than the rest. Part-time jobs for similar training work pay $100-200/hour and highly-experienced software engineers are hired by the dozen – many of whom work at Big Tech! Here's a job listing from G2i:

> "We're looking for highly experienced software engineer to help evaluate the quality of interactions with modern coding agents such as OpenAI Codex and Claude Code.
>
> This is not a traditional engineering role.
>
> You won't be writing production code.
>
> You'll be evaluating something harder: whether the model thinks like a great engineer."

If anything, the contracted devs are more motivated and the work is more targeted. Of course, it's also more generic and Meta can probably produce a model that's more fine-tuned to its own internal codebase with their approach.

OpenAI starts moving off Datadog after record bill?

In October, I reported that OpenAI reportedly set a new record spend on Datadog of $170M per annum (!!), beating previous record holder Coinbase's $65M annual bill.

At the time, Coinbase quickly took action by building an internal observability platform, and planned to move off. Eventually, Datadog managed to keep its customer with a steep discount.

It was predictable that OpenAI would not stomach astronomical observability fees either: in October I mused on whether the Statsig acquisition might result in OpenAI building its own stack. As per a recent JP Morgan equity report, OpenAI is moving log management – a subset of Datadog's functionality – into ClickHouse.

I assume OpenAI will keep looking for cheaper alternatives for the other parts of its stack as well, and that Datadog will probably throw massive discounts at the AI company to retain it.

Apple lets slip that it uses Claude

Last week, in an update to the Apple Support app, a Claude.md file was accidentally shipped, Mac Rumors analyst Aaron Perris shared. This file reveals that engineers are using Claude Code to build the app, as well as some internal rules for development:

- Use AsyncStream when showing realtime messages, instead of the Combine framework that the rest of the app uses
- Use Swift's actors for thread-safe concurrent messaging, and do not mark threads as @MainActor (that would force everything onto the main thread)
- Instructions on conditional compilation with flags like #if JUNO_ENABLED, #if DEV_BUILD
- Store sensitive session info in Keychain
- Instructions on how to route messages based on roles (.client / .agent / .assistant)

Apple issued an emergency app update to remove the file, which also exposed internal conventions on the dev team for the tool. *It's a good reminder to check that internal files are not shipped into production.*

Executive transfers at Microsoft

Big changes are taking place at Microsoft, with four executives from GitHub and CoreAI (CoreAI is the division that owns GitHub) moving to the Xbox division:

- Jared Palmer: a senior vice president (VP) at GitHub, previously working on building stacked diffs into the platform. *We previously covered more on this technique in the deepdive Stacked diffs (and why you should know about them)*
- Tim Allen: GitHub's senior VP of design and research

- Jonathan McKay: Head of Growth at CoreAi
- Evan Chaki: a general manager at CoreAI

The strange thing is Microsoft's CoreAI division *should* be where future growth will be! CoreAI includes Copilot, GitHub, and the AI stack. And yet, very senior people seem to be deflecting to Xbox; a division where growth is slowing and the total market size is likely smaller than the AI market.

I asked around internally about what's going on, and kept hearing the same thing: *leadership*. The new leader of Xbox, Asha Sharma, has a reputation as an incredible leader. She previously worked at CoreAI, and some of those moving to Xbox are former colleagues, and others she convinced to join.

Of course, when people move jobs for a leader, the unspoken bit is that they probably don't rate their present leaders too highly, or don't see comparable career options from staying put.

I've also heard that more people from CoreAI are considering moving over to Xbox. To make it more confusing, Xbox has decided to retire all its Copilot products; not all that surprising, given gamers passionately hated Copilot for Xbox. It's almost like Microsoft had two organizations working against each other: pro-Copilot (CoreAI) and anti-Copilot (Xbox).

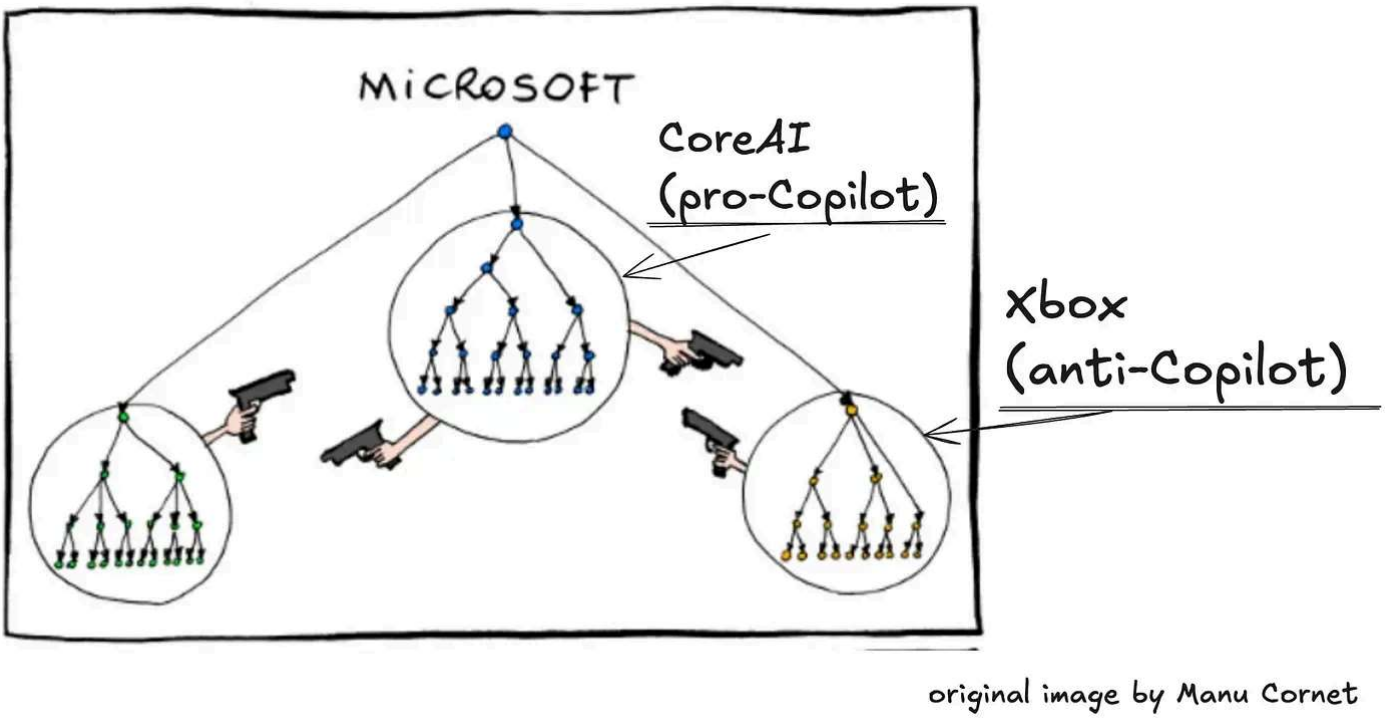

Updated mental model of Microsoft, the original from tech cartoonist Manu Cornet

VS Code inserted "coauthored by Copilot" regardless of usage

Well, this was nasty: three weeks ago, Microsoft made a change to VS Code where it started to insert "coauthored by Copilot" to the commit history of any and all changes made by VS Code with Copilot enabled – whether or not Copilot was actually used. This included when a developer deleted the Copilot changes, so that no change was *actually* made by Copilot.

The developer community – unsurprisingly – was outraged. Not only does this message smell like a growth hack, there is also no opt-in to this functionality, no way to turn it off, and the message is added even when Copilot's output has not been used!

The developer at Microsoft who approved the pull request apologized and shared more context:

> "I am the person who approved this PR and would like to acknowledge and apologize for the mistake of turning this feature on by default without sufficient upfront validation.
>
> There was no ill intent by evil corporation, but rather a desire to support functionality that some customers expect of VS Code w.r.t. AI-generated code. As folks mentioned here - many similar tools do this as well.
>
> Obviously, it should not be on when disableAIFeatures is on and it should not be reporting changes that were not done by AI. I'll work on fixing those and meanwhile revert default to off in 1.119 update."

We previously explored whether a tool should auto-add itself as a contributor to PRs: Claude Code does so, while Codex and OpenCode do not. My take is that this behaviour should ideally be opt-in; at the very least, there should be an opt-out option.

Coinbase layoffs: too much AI, or not smart-enough AI?

Coinbase has laid off 14% of staff, around 650 people. CEO, Brian Armstrong, shared a detailed note on the reasons:

> "First, the market. Coinbase is well-capitalized, has diversified revenue streams, and is well-positioned to weather any storm. Crypto is also on the verge of the next wave of adoption, with stablecoins, prediction markets, tokenization, and more taking off. However, our business is still volatile from quarter to quarter. While we've managed through that cyclicality many times before and come out stronger on the other side, we're currently in a down market and need to adjust our cost structure now so that we emerge from this period leaner, faster, and more efficient for our next phase of growth.
>
> Second, AI is changing how we work. Over the past year, I've watched engineers use AI to ship in days what used to take a team weeks. Non-technical teams are now shipping production code and many of our workflows are being automated. The pace of what's possible with a small, focused team has changed dramatically, and it's accelerating every day."

To put it simply:

1. **Coinbase's business is down ~20%** because the crypto market is in a downturn, and Coinbase takes a cut from crypto transactions. Specifically, revenue from a year ago reduced by around 23%, and the company made a loss in the last quarter (negative net income.)
2. **AI efficiencies observed.** An observation that teams are more productive with AI.

The interesting thing is that the layoffs make business sense just based on #1: the fact that Coinbase's business is struggling. But even though Amstrong says AI is making people more productive, this productivity does not seem to create revenue growth. Business seems to depend on crypto market behavior, with nothing to do with AI-enhanced productivity – or lack thereof.

Armstrong also announced other big changes:

- **Fewer managers, fewer layers, 15+ direct reports.** The structure of the company will be a maximum of five layers from the CEO, with some managers having 15+ direct reports. This flattening is similar enough to what Meta did back in 2023, and the trend is playing out again in 2026.
- **Every manager now also does IC work.** The new expectation is for all managers to be "player coaches." While this sounds good in theory, it typically means managers either stop paying attention to the "people" side of their job, or do an

equally bad job at both the IC and management work. X (formerly Twitter) made this change after Elon Musk took over. The business still runs, but revenue is dropping, and it's losing market share to rival social media like Threads.

- **"AI-native pods:"** again, exactly the same as Meta started in February, about which I shared details on, back then. Coinbase also plans to have one-person teams, where a dev plays the PM and designer role as well.

I wonder if any of these changes will boost revenue, or will they cut costs? My sense is that Coinbase is going for cost reduction and hopes to do the same work with fewer people and more AI tools.

The hard question is beyond cost reduction, can these AI tools generate more revenue? For as long as Coinbase's business rests on taking a cut from crypto transactions, the answer will most likely be "no." This suggests that whenever the crypto market enters a "crypto winter", Coinbase will probably do layoffs irrespective of AI.

72 Likes · 2 Restacks

Discussion about this post

Comments Restacks

Write a comment...

Dr Milan Milanović Tech World With Milan Newsletter 5d

Liked by Gergely Orosz

"Meta and Amazon's CEOs say teams of 5-10 devs do better work than teams of 50." This was always the case, not only with AI. Brook's Law (https://lawsofsoftwareengineering.com/laws/brooks-law/) and The Ringelmann Effect affect this (https://lawsofsoftwareengineering.com/laws/ringelmann-effect/).

LIKE (12) REPLY SHARE

Matt Kornfield Matt Kornfield 3d

On DOAC this women was interviewed about how all these jobs (not just software) will all become data labelling for AI. Pretty interesting to see it playing out at Meta.
https://www.linkedin.com/posts/karendhao_theres-a-reason-why-diary-of-a-ceo-has-built-ugcPost-

7442915398104223744-JMwe?utm_source=share&utm_medium=member_desktop&rcm=ACoAAAwIE1ABKr6Zh-boZ1gFYtJLnLhdsaE1XiU

LIKE REPLY SHARE

2 more comments...

© 2026 Gergely Orosz · Privacy · Terms · Collection notice
Substack is the home for great culture