THE PULSE

The Pulse: ‘Tokenmaxxing’ as a weird new trend

... which will probably be the shortest-lived trend because it’s so wasteful. Also: coding AI agent subsidies could be ending, Cal.com going closed source and blaming it on AI, and more.

GERGELY OROSZ
APR 16, 2026 · PAID

78 5 1 Share

The Pulse is a series covering events, insights, and trends within Big Tech and startups. Notice an interesting event or trend? Hit reply and share it with me.

Today, we cover:

1. **Tokenmaxxing: weird new trend.** At Meta, Microsoft, Salesforce and other large companies, devs are purposefully burning tokens (and money!) to inflate their AI usage and hit AI usage metrics which they treat as targets.
2. **Are coding AI-agent subsidies doomed?** At the same time as Anthropic stopped subsidizing enterprise plans, Uber managed to burn through its entire 2026 AI token budget in just 3 months. I expect per-engineer AI budgets to be rolled out across more companies soon.
3. **Industry Pulse.** The myth of Claude Mythos, Claude’s degradation, Cal.com going open source due to AI threat, Vercel open sources its “agent factories” tool, sensible AI usage guidelines in the Linux kernel, and more.
4. **Cal.com goes closed source – but is it really because of AI?** The open source Calendly alternative moved a good part of its code to a closed repo, citing AI and security concerns. But perhaps this was just a business model change that would have happened, AI or not.

1. Tokenmaxxing: weird new trend

Inside Meta, an engineer created a “token leaderboard” that ranks employees by token usage. Last week, The Information reported:

> “Employees at Meta Platforms who want to show off their AI superuser chops are competing on an internal leaderboard for status as a “Session Immortal”— or, even

> better, "Token Legend."
>
> The rankings, set up by a Meta employee on its intranet using company data, measure how many tokens — the units of data processed by AI models — employees are burning through. Dubbed "Claudeonomics" after the flagship product of AI startup Anthropic, the leaderboard aggregates AI usage from more than 85,000 Meta employees, listing the top 250 power users.
>
> The practice is emblematic of Silicon Valley's newest form of conspicuous consumption, known as "tokenmaxxing," which has turned token usage into a benchmark for productivity and a competitive measure of who is most AI native. Workers are maximizing their prompts, coding sessions and the number of agents working in parallel to climb internal rankings at Meta and other companies and demonstrate their value as AI automates functions such as coding."

I spoke with a few engineers at Meta about what's happening, and this is what they said:

- **Massive waste.** Plenty of devs are running an OpenClaw-like internal agent that burns massive amounts of tokens for little to no outcome.
- **Outages caused by AI overuse.** A dev mentioned that some SEVs were caused by what looked like careless AI code generation; almost like a dev behind the SEV was more concerned with churning out massive amounts of code with AI than with product quality.
- **Gamified leaderboard.** Those at the top of the leaderboard produce throwaway, wasteful work. This is painfully clear to anyone who checks Trajectories (AI prompts), which can be viewed.

As per The Information, Meta employees used a total of 60.2 trillion AI tokens (!!) in 30 days. If this was charged at Anthropic's API prices, it would cost $900M. Of course, Meta is likely purchasing tokens at a discount, but that could still come in at $100M+ – in large part from senseless "tokenmaxxing".

After backlash on social media, Meta abolished the internal leaderboard last week. One day after The Information revealed details about the incredible tokenmaxxing numbers, I confirmed that Meta has taken down its leaderboard; perhaps they realized that the incentive created enormous and unnecessary waste. If

so, it's a bit surprising that it took media coverage for the social media giant to reach that conclusion.

One engineer at Meta told me they think Meta had a different goal with the token leaderboard. A long-tenured engineer suspects increasing AI usage actually was the real goal. They said:

> "Putting a leaderboard in place was always going to incentivize much more AI usage. And more AI usage means producing a lot more real-world traces. These traces can then be used to train Meta's next-generation coding model better.
>
> I believe this was the goal, even if no one said it out loud.
>
> It's an expensive way to generate data for training, but if any company has the means to do so, it's Meta."

Microsoft: full-force tokenmaxxing

Similarly, Microsoft has had an internal token leaderboard like Meta's since January, and it started pretty well, as I reported back at the time: there's an internal token dashboard that displays the individuals who use the most tokens in order to promote the use of tokens and experimentation with LLMs. At the Windows maker, this leaderboard is interesting:

- Very senior engineers – distinguished-level folks – are in the top 5 across the whole company, despite the fact that this group generally wrote little code in the past.
- VP-level folks make the top 10 and top 20, despite often being in meetings for most of the day and rarely writing code.

However, what starts as a metric for performance reviews or promotions can quickly become a target for devs. I talked with a software engineer at the Windows maker who admitted they're full-on "tokenmaxxing" – not to get on the leaderboard, but rather because they don't want to be seen as using too few tokens:

> "We have internal dashboards and metrics tracking AI usage, token usage, percentage of code written by AI vs hand-written code.
>
> I am conscious of not wanting to be seen as "uses too little AI," and I'm not ashamed to say I need to do tokenmaxxing to do this. Things I do to inflate my

token usage metrics:

- Ask AI questions about the code already in the documentation. The AI pulls up the documentation, processes it, and gives me results 10x slower, but while burning lots of tokens. I could use "readthedocs" [an internal product], but then my token numbers would be lower
- Ask the AI to prototype a feature that I have no intention of working on. Prompt it a few more times, then throw the whole thing away
- Default to always using the agent, even when I know I could do the work by hand much faster. Then watch it fail"

This engineer is relatively new at the company, so is concerned about job security, and is playing this game to avoid being tagged as insufficiently "AI-native" by burning far more tokens than necessary.

Salesforce: burning tokens to hit "minimum" & "ideal" targets

Elsewhere, Salesforce has created "tokenmaxxing" incentives, as well. Talking with an engineer there, I learned that the company built two tools that effectively incentivize excessive spending on tokens:

1. **"Minimum" incentives with a tracking tool.** There's a Mac widget that shows your own spend, updated every 15 minutes. It also displays minimum expected spend. Last week, the target was $100 on Claude Code, and $70 on Cursor.
2. **Showing everyone's spend.** A web-based tool to see the token spend of any colleague. It's used to check where team mates' usage is at.
3. **"Maximum" spend limits that can be exceeded.** Up to a week ago, there was also a *maximum* monthly limit of $250 for Claude Code and $170 for Cursor. *However, this can be exceeded with the simple press of a button if the limit is reached. I've learned that last week, some engineering organisations at Salesforce had their "maximum" limit removed in order to "remove any friction from the development process."*

The message Salesforce sends to staff is clear: "use a minimum of $170/month tokens or be flagged." Who wants to get flagged for using too few tokens? The outcome is somewhat wasteful token spend:

- **Burning tokens for nothing.** Devs ask Claude or Cursor: "build me X," where X is a project or product with nothing to do with their work, and not something they'd ever ship. It's just a way to burn tokens
- **Calibrating token spend to be above average.** Plenty of devs browse peers' token spend to figure out the slightly-above average point, then use the tokens needed to hit that mark

Shopify: an example on how to avoid tokenmaxxing

The first-ever token leaderboard that I'm aware of was built by Shopify in 2025. And it worked well! Last June, the Head of Engineering at Shopify, Farhan Thawar, told me on The Pragmatic Engineer Podcast:

> "We have a leaderboard where we actively celebrate the people who use the most tokens because we want to make sure they are [celebrated] if they're doing great work with AI.
>
> [And for the top people on the leaderboard,] I want to see why they spent say $1,000 a month in credits for Cursor. Maybe that's because they're building something great and they have an agent workforce underneath them!"

I asked Farhan for details on how it's gone since. Here's what he told me:

> "We have since renamed the token leaderboard to usage dashboard: for obvious reasons, as we don't want to encourage "competing" to make it to the top of this board. We have token spend on our internal wiki profile as well as on the usage dashboard.
>
> **We also have circuit breakers to catch "runaway agents."** So if personal spend spikes within a day, we can cut off access immediately, and you can renew if the usage spike was deliberate, or if it was a runaway agent. The circuit breaker worked well for us: we've not only caught runaway agents, but found bugs in our infra this way!"

Shopify's approach seems to have worked for a few reasons:

- **The usage dashboard served as a "push" for devs to use AI tools, early-on.** Last year, devs were mostly experimenting with AI tools because they were not as

performant as today. The usage dashboard encouraged developers to try new tools, and highlighted power users.

- **Circuit breakers helped.** Cutting off spend when usage spikes helped catch "runaway agents."
- **High usage is looked at.** Farhan checks-in with top-spending individuals to understand the use cases. Any tokenmaxxing would likely have been spotted at this stage, which would have been a bit embarrassing for the user!

One more interesting learning Farhan shared with me: it's more interesting to not look at "who spent the most in *overall* token cost?" but instead, "whose *tokens* cost the most?" Devs who generate tokens that come out as expensive have turned out to do in-depth work that was interesting to learn about!

Tokenmaxxing: great for AI vendors, bad for everyone else

I see very few rational reasons why incentivizing tokenmaxxing makes sense for any company. It results in increasing AI spend – by a lot! – in return for little to no value. Heck, in some cases it actually incentivises slower work – as shown by devs using the AI to answer questions when documentation is readily available – and encouraging 'busywork' where devs prompt projects that they don't even want to ship. Tokenmaxxing seems to push devs to focus on stuff that makes no difference to a business.

It feels to me that a good part of the industry is using token count numbers similarly to how the lines-of-code-produced metric was used years ago. There was a time when the number of lines written daily or monthly was an important metric in programmer productivity, until it became clear that it's a terrible thing to focus on. A lines-of-code metric can easily be gamed by writing boilerplate or throwaway code. Also, the best developers are not necessarily those who write the most code; they're the ones who solve hard problems for the business quickly and reliably with – or without – code!

Similarly, the number of tokens a dev generates can easily be gamed, and if this metric is measured then devs will indeed game it. But doing so generates a massive accompanying AI bill!

2. Are coding AI-agent subsidies doomed?

Two recent events may be related:

Uber blew through its entire 2026 AI budget in only 3 months. As per Uber CTO, Neppalli Naga, in The Information:

> "I'm back to the drawing board because the budget I thought I would need is blown away already."

From my experience of working at Uber, the company is pretty good at forecasting costs for services and tools. Spending the annual AI budget for the year in 12 weeks suggests the true yearly budget could easily be 4x higher or more than planned. This usage seems to be coming from Claude Code which is extremely popular in the ridesharer's engineering department, where Cursor's popularity is declining.

Elsewhere, Anthropic removed subsidized usage from Enterprise plans. Anthropic updated its pricing documentation last week, making this point:

> "[Enterprise plans follow] usage-based pricing. The seat fee only covers access to the platform and doesn't include any usage. All usage across Claude, Claude Code, and Cowork is billed separately at standard API rates, based on what your team actually consumes. There are no per-seat usage limits and no included token allowance.
>
> Admins can set spend limits at the organization and individual user levels to manage costs."

This means that only individual customers get increased usage included in Claude Max plans! As a result, Anthropic is no longer offering subsidies for enterprise customers.

I suspect that Uber eating up its AI budget so fast due to Claude Code's price rise, and Anthropic moving Enterprise customers to token-based billing, may be two separate events that point to the same thing.

I expect large companies to start setting per-engineer limits on AI usage, and look to cheaper models as well. The cost of giving access to state-of-the-art (SOTA) models like Opus 4.6 and GPT-5.4 is becoming ever more costly, even for the largest tech companies. A couple of options they have:

- Set limits on SOTA model spend to incentivize using the most capable models only for the most complex tasks

- Offer less capable, much cheaper models for "everyday" coding tasks. These might be inexpensive third-party models, or even hosting AI models on their own infrastructure!

For leading AI vendors, a push back by companies against paying more for AI usage could be bad news. It looks likely that the "tokenmaxxing" trend will be one of the shortest-lived we've seen.

3. Industry Pulse

Claude Mythos: myths about a supposedly good model

Last week, Anthropic announced its latest, most powerful model, Mythos. Unlike past models, this one has not been released because Anthropic claims it's so *powerful* and has such strong cybersecurity capabilities that it can find zero-day vulnerabilities. So, Anthropic is launching a project called Glasswing, where they give access to Mythos to select companies like Amazon, Apple, Microsoft and others, in order to secure their own software and open source software first.

Anthropic released a lengthy research paper about Mythos, but it's hard to tell how much of this is marketing hype. Is Anthropic holding back this model because it's dangerously powerful, or to generate market buzz about it, or is Mythos simply too expensive to run for the capabilities it has?

We simply don't know. I talked with a current Anthropic engineer who has been using Mythos, and claims to be blown away by it. Then again, I find it hard to have an opinion about a model that cannot be used, and about which the only information comes from a vendor incentivized to talk up its own model as much as possible!

I sense that Anthropic is building their own myth, and I refuse to believe that the company has created a model that other labs could not replicate in a month or two. I'm also reminded of events in 2019 when OpenAI withheld GPT-2 from the public. Here's some headlines about that from the time:

- "OpenAI builds artificial intelligence so powerful that it must be kept locked up for the good of humanity" – Metro
- "Our text generator is so good it's scary." – CNET
- "AI can write just like me. Brace for the robot apocalypse" – Guardian

I remain skeptical that a model is "too good" to be released, and my hunch is that there are mundane reasons behind Anthropic shipping a PDF instead of a model.

Claude seems to be degrading, but no one knows why

There are lots of complaints online about how Claude seems to have become "dumber" and "lazier," I have also noticed this in my own experience.

Director of AI at AMS, Stella Laurenzo, published a long, detailed issue with Claude Code describing exactly this:

> "Claude has regressed to the point that it cannot be trusted to perform complex engineering.
>
> We have a very consistent, high complexity work environment and data mined months of logs to understand why – essentially – starting in February, we have noticed a degradation performing complex engineering tasks. Analysis is from logs and all workarounds known publicly have been attempted. Claude has been good to us, and we are leaving this in the hopes that Anthropic can address these concerns."

The issue has detailed logs, and documents the degradation of model output. Coincidentally, the degradation started around when Anthropic supposedly rolled out Claude Mythos to its engineering team, internally (February 24).

Could it be that Mythos has such high resource usage that Anthropic has been finding ways to allocate less compute to public Claude models? Or could Anthropic be allocating less compute to Claude for other reasons?

The problem is that a closed model like Claude offers no transparency about what is happening, and why! Therefore, Anthropic customers get frustrated but get no answers, and are left with a worse model than a few months ago.

I would expect larger companies to start to consider self-hosting models on their own infrastructure. That way, they can at least guarantee that if a model works a certain way one month, it will work the same way the following month. Using a third-party service is looking like a risk, especially with Claude degradations.

Vercel open sources "software factories" tool

A new and interesting open source project coming from Vercel is open-agents.dev. Founder and CEO, Guillermo Rauch, explains:

> "You've heard that companies like Stripe (Minions), Ramp (Inspect), Spotify (Honk), Block (Goose), and others are building their own "AI software factories". Why?
>
> 1. On a technical level, off-the-shelf coding agents don't perform well with huge monorepos, don't have your institutional knowledge, integrations, and custom workflows.
> 2. On a business level, the moat of software companies will shift from 'the code they wrote', to the 'means of production' of that code. The alpha is in your factory.
>
> Open Agents deploys to Vercel's agentic infrastructure:
>
> Fluid for running the agent's brain, Workflow for its long-running durability, Sandbox for secure code execution, AI Gateway for multi-model tokens."

Many tech companies are indeed building their own "background agents" infrastructure. Vercel is making something that resembles *some of this,* but it necessarily lacks the need to be integrated into a monorepo. This project *could* make it easier to build such a platform. However, it seems as if a bunch of extra work would be necessary to integrate this platform with internal codebases.

I find it interesting that even vendors like Vercel realize that it will be a big enough business opportunity to run internal coding agents at various tech companies.

Meta on track to become a bigger ads business than Google

The Wall Street Journal reports that Meta is on track to generate more ads revenue than Google this year, thanks to Meta's ads business growing around 24%, while Google's growth is around 11%.

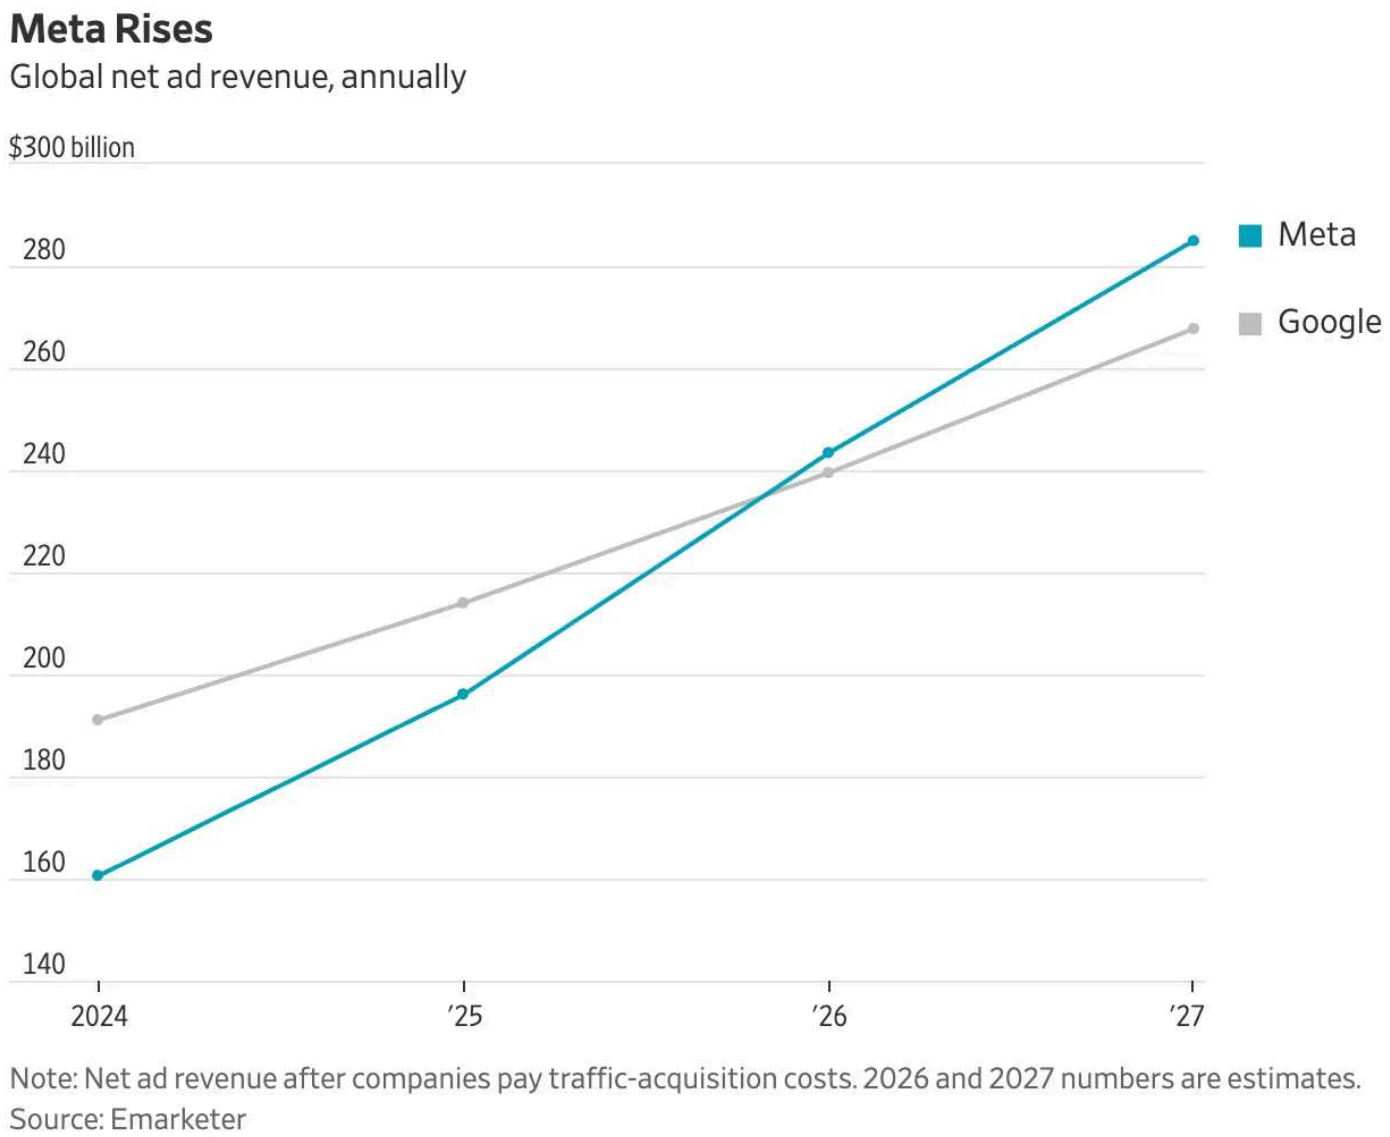

Projected annual revenue from ads, in billions. Source: The Wall Street Journal

I had assumed that Google would always be the world's largest advertising business, so for a social media company to overtake it indicates that Google made a big mistake by missing out on social media. For Meta, investing heavily in ML and AI systems for ads targeting on Facebook and Instagram is paying off handsomely.

Meta's revenue growth is even more impressive given how its ads growth flatlined fairly recently, in 2021-2022. But the company had pulled off a remarkable turnaround by 2024:

> "Less than a year ago, things looked very grim at Meta. In November 2022, it hit a 7-year low in stock price, revenue growth stalled, and the company did mass layoffs by firing 13% of staff (11,000 people) in what was its first-ever mass cull of jobs. But it wasn't over, just 5 months later, Meta announced another 10,000 jobs were being cut. In six months, Meta let go 1 in 4 employees, many of them software engineers.

> Just as quickly and unexpectedly, an incredibly fast business recovery is underway. The company's stock and valuation are at an all-time high of $1.1T at time of publication. This is 4.5x of what it was in November 2022!"

While Meta is on course to be a bigger ads business than Google this year, the search giant possesses far more business units than Meta:

- Cloud unit: Google Cloud
- Mobile phone platform and hardware platform: Android and Pixel phones
- AI business: selling Gemini models
- Business software suite: Google Workspace
- Consumer software: Gmail (the leading email platform) and Google Maps
- Self-driving division: Waymo
- ... and many smaller business divisions

Today, Google is valued at about twice that of Meta (Google: $4.0T, Meta: $1.7T). I'm assuming this gap is because Google has a number of business areas to grow in the future, aside from only ads.

Sensible AI usage guidelines in the Linux kernel project

Linus Torvalds published AI contribution guidelines for the project which are worth reporting because they are so refreshingly clear and sensible:

- AI tools should follow the standard kernel development process
- All code needs to be compatible with the Linux kernel license (GPL-2.0)

And the most important part:

> "**Only humans can legally certify the Developer Certificate of Origin (DCO).** The human submitter is responsible for:

- Reviewing all AI-generated code
- Ensuring compliance with licensing requirements
- Adding their own signed-off-by tag to certify the DCO
- Taking full responsibility for the contribution"

The basic message is to use as much AI as you want, in the way you want, but in the end you need to be fully responsible for any code that you used AI for!

This is straightforward and easy, and makes me wonder why isn't AI usage governed this way for all projects, and inside every company?

Relatedly, we previously covered How Linux is built with longtime Linux maintainer, Greg-Kroah Hartman.

4. Cal.com goes closed source – but is it really because of AI?

Cal.com is a popular open-source Calendly alternative. At least, it *used to be* open source because this week, the company announced they are moving to a closed source model. The reason? The threat of AI – specifically, the threat of AI being used to detect security vulnerabilities. Cofounder and CEO, Bailey Pumfleet, said:

> "In the past, exploiting an application required a highly skilled hacker with years of experience and a significant investment of time to find and exploit vulnerabilities. The reality is that humans don't have the time, attention, or patience to find everything.
>
> Today, AI can be pointed at an open source codebase and systematically scan it for vulnerabilities.
>
> Being open source is increasingly like giving attackers the blueprints to the vault. When the structure is fully visible, it becomes much easier to identify weaknesses and exploit them.
>
> This uncertainty forced us to make a choice: remain open source and accept increasing risk to customer data, or move to closed source to reduce that risk. It's not a perfect solution, but we have to do everything we can to protect our users."

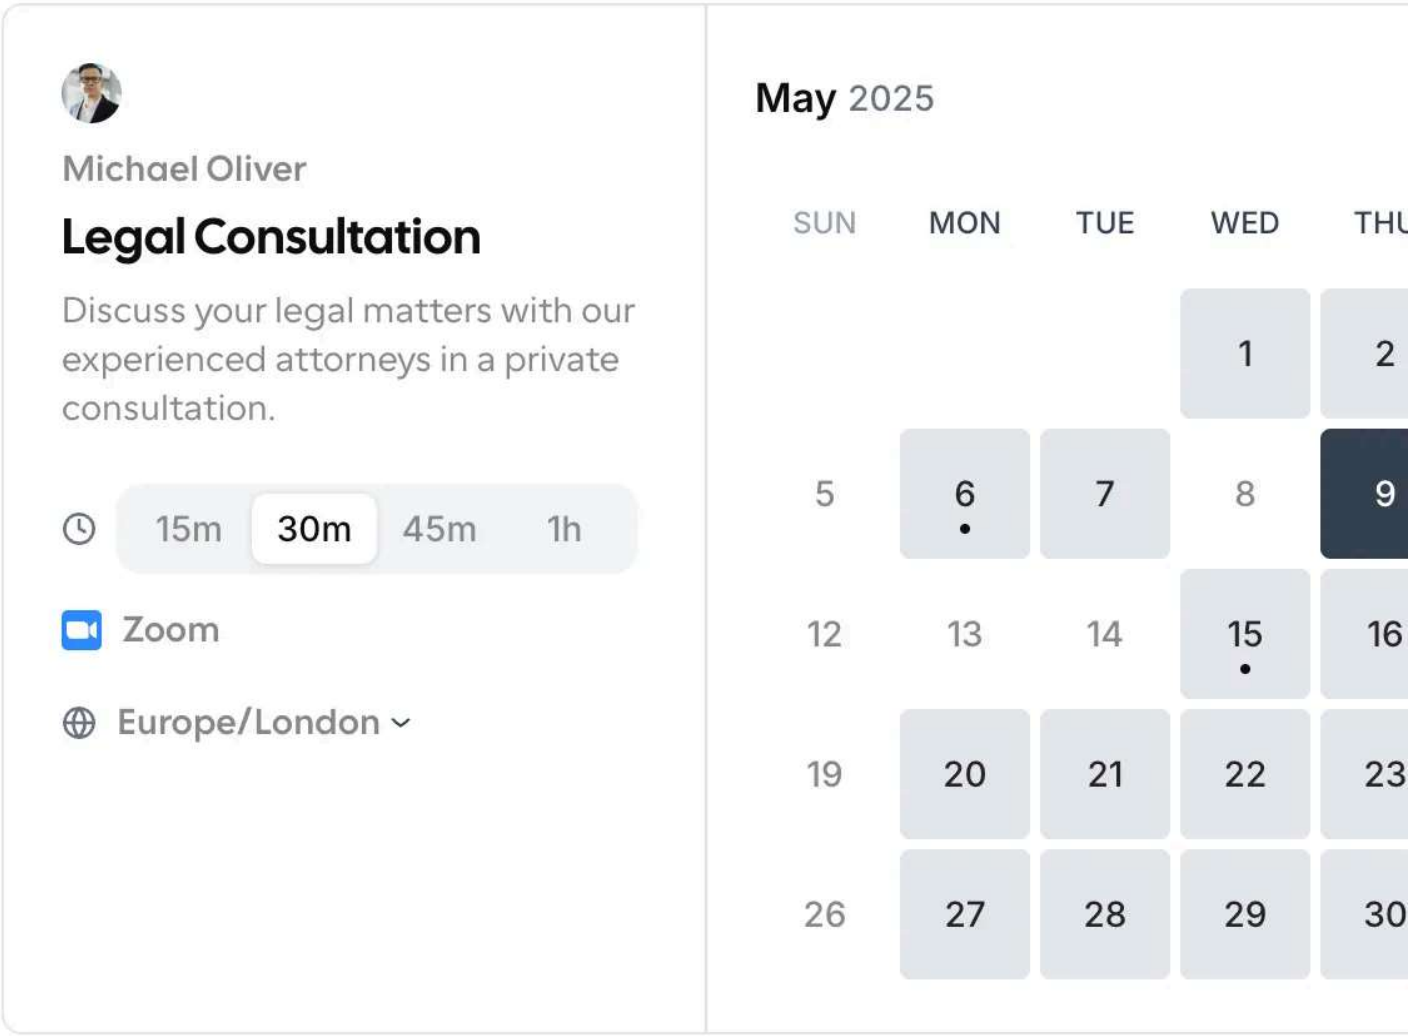

Cal.com: one of many scheduling solutions, with advanced capabilities and a lot of integrations. The main differentiation was the open source nature, until now

The company made the change to closed source in two steps:

- They relaunched basic scheduling functionality as Cal.DIY, under an MIT license – one of the most permissive licenses. This functionality is what individuals and hobbysts need.
- They moved commercial and enterprise features – like teams, ruting forms, auth, and others – into a closed-source repo.

This decision comes a week after Anthropic announced Mythos, and personally, I'm not entirely convinced that Cal.com has gone closed source over AI-security concerns. Other potential reasons:

- Being open source makes it t*oo easy* for any team to copy Cal.com's product – *especially* with AI.
- Cal.com already became the #1 Calendly alternative. Being open source helped them get to this position; but going forward, staying open source would not push

them forward

- The AI-security angle is an ideal excuse to go closed source: it makes their product more “resilient” against being copied by AI tools, and widens the product’s moat.
- Being “the open source Calendly alternative” was strong marketing, and even a good distribution strategy when launching the business. Today, Cal.com seems to be moving more to an “open core” model where the “hobbyst” part of the business remains open source, and the commercial/enterprise parts is a closed core.

My sense is that this change probably would have happened independent of AI. Going from an open source product to an open core commercial product is the most well-travelled path in commercial enterprise software. *It’s also how HashiCorp commercialized, as we previously discussed [with HashiCorp founder Mitchell Hashimoto](#).*

78 Likes · 1 Restack

Discussion about this post

Comments Restacks

Write a comment...

Matt Khor 16 kwi

Liked by Gergely Orosz

The tokenmaxxing thing is really interesting to me because it seems to intertwined with the productivity conversation. We really don't know who is just burning tokens for no reason and who is using it to do actual real work because measuring productivity in a single metric has never been and continues to not be possible.

I feel like early on in the AI boom all token usage was generally considered close to free and the general sentiment was that you should be using as many tokens as you can so that you're not being left in the dust by powerusers. Now, it feels a little more like we're entering an era where people will need to get better at efficiency/usefulness per token. I wonder if there will be meta tooling at some point to let people label prompt outcomes.

LIKE (5) REPLY SHARE

1 reply

Rizzy Smoove 8d

It's troubling to know real climate damage is underway as a result of promotion targets.

LIKE (2) REPLY SHARE

3 more comments...

© 2026 Gergely Orosz · Privacy · Terms · Collection notice
Substack is the home for great culture